



UNIVERSITÄT
LEIPZIG



Institute of Computer Science
Faculty of Mathematics and Computer Science
ScaDS.AI Dresden/Leipzig

Automated Multi-Hop Dataset Generation for Benchmarking Cited Text Span Extraction in Scientific Question Answering

Bachelor's Thesis Exposé

submitted by:
Tom Paul Gresens

Matriculation Number:
3779495

submitted on:
July 15, 2026

Supervisors:
Prof. Dr.-Ing. habil. Erik Buchmann
Prof. Dr.-Ing. Michael Färber
Tobias Schreieder

© 2026

This work, including all its parts, is **protected by copyright**. Any use beyond the strict limits of copyright law without the consent of the author is prohibited and punishable by law. This applies in particular to reproductions, translations, microfilming, and storage and processing in electronic systems.

Contents

1. Introduction	1
2. Related Work	3
2.1. Datasets	3
2.2. Cited Text Span Extraction	4
2.3. System Evaluation	5
3. Methodology and Approach	8
3.1. Dataset	8
3.1.1. Document selection	8
3.1.2. Question-Answer generation	8
3.1.3. Quality assurance	9
3.2. Cited text span extraction	10
3.2.1. Native SQuAI extraction baseline	11
3.2.2. Dense retrieval	11
3.2.3. Sparse retrieval	11
3.2.4. Hybrid retrieval	12
3.2.5. Cross-Encoder reranking	12
3.2.6. Generative LLM	12
3.3. System evaluation	13
3.3.1. Upstream system evaluation	13
3.3.2. Span evaluation	14
4. Timetable	15
Bibliography	16

1. Introduction

Large Language Models (LLMs) have fundamentally transformed how we look up information and obtain answers to questions, becoming an increasingly important tool, not only in everyday life, but also across a range of professional fields. With the use of such systems, the need to verify the factual correctness and easily do further research on the underlying sources arises, especially when answers lack important details or provide a depth of explanation that is misaligned with the user's existing knowledge. This is especially important in high-stakes fields such as medicine, law or scientific research, where seemingly plausible but unbacked or false claims could lead users to draw false conclusions.

One increasingly popular method is the approach of Retrieval Augmented Generation (RAG), replacing zero-shot LLM solutions with a two-step process. Firstly, relevant sources in the form of documents or other multi-media data are retrieved and then fed into the generating model for it to use as ground truth for answer generation and later to enable citation extraction. Retrieving correct and precise text-spans alongside the generated answer is not a trivial task, and with domain-specific requirements widely varying, such systems are part of active research [1] [2]. One system applying such a RAG approach is SQuAI [3], aimed at answering complex open-domain scientific questions on its knowledgebase containing 2.3 million full-text scientific papers. For further performance improvements, it introduces a multi-agent pipeline, enhancing document retrieval capabilities as well as providing document-level backlinks for parts of its generated answer. This system, while providing some backlinking, does still fall short of providing more granular, sentence-level cited text spans, which have proven to be much more useful to users, as they are quicker and easier to verify, enhancing trustworthiness and usability of the system.




Figure 1.1.: Sentence level cited text span attribution

Therefore, this thesis will focus on benchmarking different cited text span extraction methods to enhance source attribution, leading to the following core contributions:

1. Curate a dataset of high-quality, multi-hop scientific questions building upon a large-scale retrieval corpus that require complex reasoning capabilities in order to be answered and closely resemble real-world scientific inquiry needs.
2. Introduce a hierarchy of different cited text span extraction methods ranging from simple lexical and semantic matching up to complex reasoning capabilities, as well as strategic combinations of such, to allow for tradeoffs in accuracy and efficiency.
3. Apply an automated framework to benchmark the retrieval and answering capabilities of non-parametric, retrieval- and evidence-based text generation systems on the proposed dataset as well as evaluate the extraction methods in terms of efficiency, correctness and quality of source attribution using NLI models and LLM-as-a-judge approaches.

2. Related Work

This chapter provides a comprehensive review of prior work across three key dimensions relevant for the development of a scientific QA system capable of providing cited text spans as evidence for its generated answer. First, it examines existing scientific QA datasets with respect to their suitability for SQuAI. It then compares different established cited text span extraction paradigms and their tradeoffs in terms of precision, efficiency and computational requirements before discussing automatic approaches for evaluating text span attribution quality in long-form, multi-hop settings.

2.1. Datasets

Utilizing SQuAI as a system basis, this work conducts a systematic evaluation of cited text span extraction methods for fine-grained source attribution in scientific QA. Therefore, a highly specialized dataset of open-domain, long form scientific questions is needed. Due to the scarcity and high cost of curating such datasets, only a limited number of suitable options are available. With SQuAI being designed as a system to answer complex scientific inquiries based on the unarXive 2022 Dataset [4], further requirements are introduced, namely:

1. **Domain Alignment:** Questions must fall under the knowledge scope of the underlying knowledgebase, in order for the system to be able to provide cited text spans as proof alongside its answer. In terms of the SQuAI System, this boils down to questions targeting physics, mathematics and computer science that are also part of the papers in the unarXive 2022 Dataset.
2. **High Complexity:** Questions must go beyond simple factoid questions and require complex reasoning capabilities to answer, as questions posed by scientists usually require synthesizing, interpreting and evaluating data, requiring high cognitive abilities to interact with. In order to be able to test if the SQuAI system can be used under such circumstances, the questions in the dataset must also have a similar degree of complexity.
3. **Multi-Hop:** Questions must require the information in multiple documents to be answered, in order to test the ability of the system to synthesize information from multiple sources, as is required in naturally occurring scientific research.
4. **Decontextualization:** Questions must be answerable in an open-domain setting and cannot contain any kind of reference to specific papers and their contents that lose meaning or become ambiguous outside of the context of their paper. As a result, questions must contain enough context such that they are standalone answerable.
5. **Text format:** Questions can only target information in a paper that is present in the form of text, as current versions of the SQuAI system do solely rely on textual information, not taking into account images, graphs or similar.

6. Ground Truth: In order to evaluate the correctness of the model’s answer, a dataset must provide a ground answer alongside each question.

However, comparing some of the most prominent datasets shows that none fits all discussed criteria, finally resulting in the need for a newly curated dataset.

Dataset	Domain Alignment	High Complexity	Multi-Hop	Decontextualization	Text format	Ground Truth	Sufficient Dataset size (>500)
Qasper [5]	partial	shallow	no	no	yes	yes	yes
ALCE [6]	no	shallow	yes	yes	yes	yes	yes
SPIQA [7]	yes	complex	no	partially	no	yes	yes
QASA [8]	yes	complex	no	no	yes	yes	yes
ExpertQA [9]	yes	very complex	yes	yes	yes	yes	no ¹
SciDQA [10]	yes	very complex	no	no	no	yes	yes
PeerQA [11]	yes	very complex	no	no	yes	yes.	yes

Table 2.1.: Comparison of Scientific QA Datasets against SQuAI Requirements

Abbreviations: CS = Computer Science; ML = Machine Learning; NLP = Natural Language Processing.

2.2. Cited Text Span Extraction

Context extraction methods can be most broadly divided into pre- and post-hoc extraction methods, depending on the time the context attribution happens. While pre-hoc systems such as QASPER [5] and QASA [8] extract citable passages before and during runtime of the question-answering system, post-hoc systems search and attach supporting evidence afterwards. Post-hoc systems have been shown to outperform pre-hoc systems, achieving higher citation coverage and answer correctness as well as lowering model hallucination [12]. While a number of systems rely on zero-shot LLM solutions for context extraction [13] [14] [15], work on the DECOMPTUNE [16] system has shown that decomposing facts and decontextualizing answer sentences before attributing source text spans with LLM prompting can yield significant improvements, especially when applied to answers for long-form, multi-hop questions. In order to eliminate the probability of hallucinated model output, models can be finetuned [17] or the answer space restricted to the start and end token of the cited text span, forcing strict citation [17].

So-called white-box methods rely on access to the model’s internal weights, rendering them incompatible with proprietary LLMs. The LAQuer [18] system, as well as earlier token-embedding approaches it is built upon [19], analyzes internal hidden states to extract the most likely text span that a user-highlighted part of the LLM’s answer is based upon. To pinpoint a certain text span, the system calculates the maximum cosine similarity between individual tokens in the source text and the averaged hidden representation of the highlighted span, for which attribution is sought. The system then extracts the precise quote from the surrounding window of text. However, this method has been shown by the authors to perform significantly worse than LLM Prompting, with drastically lower accuracy and missing suitable citations, especially in larger search spaces, e.g.

¹question-answer pairs are broadly distributed across diverse fields of study, resulting in insufficient topical overlap with unarXive 2022

longer text. Similarly, AttnAttribute [20] identifies "Attention Heads", tracking the model's internal attention values during generation time. Just like LAQuer, it then extracts text spans from the passage surrounding the anchor with the highest model attention. Benchmarks made as part of the development of DECOMPTUNE [16] concluded that both of those methods performed significantly worse than LLM Prompting.

Recent advancements have explored generative approaches. Rather than retrieving and reranking text spans, such systems generate them. In order to prevent hallucinations, the output vocabulary of the systems is typically restricted to the underlying retrieval corpus by applying constrained decoding mechanisms, such as prefix trees or FM-indices [21]. GERE [22], as one such system using a prefix tree, applies this concept of restricted output in the document retrieval step, generating document titles for selecting appropriate source texts, as well as sentence identifiers mapped to individual sentences. This system, while removing the need for a massive document index as well as showing significant improvements in precise and correct text span extraction compared to baselines such as Dense Passage Retrieval (DPR) and BM25, requires the prefix tree to be precomputed and stored in memory, making it highly inflexible for changes in the underlying corpus. Similarly, MRR-FV [23] further improves this approach with multi-hop reasoning capabilities. Instead of generating sentence identifiers, it generates sentences directly, leveraging the natural language understanding capabilities of LLMs, leading to improvements in precision and quality of attributed spans, but also further bloating the prefix tree, making the system even harder to scale.

Despite the diversity of proposed paradigms, current work mostly evaluates them in isolation, restricting comparison to baselines or direct predecessors on heterogeneous tasks, making broad comparisons for practical applicability unfeasible. Consequently, this work introduces a hierarchy of extraction methods, increasing in both reasoning capabilities and computational requirements, benchmarking both accuracy and efficiency on a unified framework, allowing for systematic trade-off analysis between paradigms in the context of open-domain multi-hop scientific QA.

2.3. System Evaluation

While human annotation has always been the gold standard for assessing natural language generation (NLG) tasks, relying solely on it is a bottleneck for developing models of increasing size and complexity. While providing feedback with extensive reasoning and natural language understanding, human evaluation suffers from being costly, time-intensive and, depending on the specific task, highly subjective. Thus, lowering or even replacing the need for human evaluation completely is an important problem.

Early metrics heavily rely on measuring the lexical overlap between text spans, for instance of a human-written ground truth and an LLM-generated answer for the same questions. The most prominent examples include BLEU [24] and ROUGE [25], both being n-gram based metrics. While the former is precision-oriented, measuring how many of its n-grams also appear in the human reference span, the latter applies a recall-oriented approach, measuring how many n-grams of the human reference are also present in the LLM-generated text span. Being computationally inexpensive and language independent, these metrics have established themselves as rigid industry standards and

baselines for other methods. Their alignment with human judgement, however, has been shown to be low, suffering from the lack of semantic understanding, leading to penalizing valid paraphrasing as well as missing contradictory or hallucinated content.

Partly overcoming those limitations, embedding-based metrics like BERTScore [26] and MoverScore [27] emerged, measuring semantic instead of lexical similarity. BERTScore does one-to-one greedy matching on vector embeddings, providing precision, recall and F1 score metrics, just like BLEU and ROUGE, but on an embedding level. MoverScore improves upon this idea, utilizing Earth Mover’s Distance (EMD) to allow for introducing one-to-many relationships, allowing for matching of single words with phrases, so-called soft matching, as well as introducing Inverse Document Frequency (IDF) for enhanced lexical matching, to better reflect the importance of specific keywords to be included. Though significantly more expensive to compute, those metrics show significant improvements in semantic robustness as well as intrinsic context awareness, due to their vector embeddings. However, both of those metrics still require a reference answer to compare to. BartScore [28] eliminates this requirement, measuring the conditional probability that one text span is generated from another using an encoder-decoder model, arguing that the probability is higher if the claims in the answer are also deeply rooted in the source text.

Besides similarity-based metrics, other research suggests approaching the task from a logical entailment perspective. The AIS [29] framework proposed a unified conceptual definition for such, arguing that a statement is considered Attributable to Identified Sources (AIS) iff completely corroborated by a source. In order to verify this, the authors imposed additional constraints being 1) that the claim made in an answer sentence must be standalone interpretable, thus decontextualized and 2) the information in the text span must support all the information present in the answer span it is attributed to. Previously discussed similarity-based metrics like BLEU, ROUGE or token-level F1, as well as solely relying on an out-of-the-box Natural Language Inference (NLI) model, have been shown to have low correlation with human judgement for both factual consistency and logical entailment assessment [30] [31]. To address these limitations, systems relying on NLI models like FactCC [31] argued that evaluating logical entailment on sentence-to-sentence level is insufficient, since information in sources might be spread across multiple passages. As a result, these systems increased their evaluation scope to the whole source text, leading to better judgement capabilities, but also increase in model input. The authors of FactCC observed that a majority of misclassifications were commonsense mistakes; the TRUE benchmark [30] observed improvements with increasing model size but also degradation with longer input, which is further supported by research on the SummaC [32] framework. To shrink input size, SummaC increased evaluation granularity by dividing both the source text and the model answer into smaller units, checking each combination of them for entailment and combining the obtained scores.

The Q² system [33] proposed using a Question-Generation - Question-Answering (QG-QA) method, generating questions to which the LLM-generated answer is the solution. Those questions are then answered with the original source document and lastly the so obtained answers are then compared with the original LLM answer using an NLI model. The authors presented significant improvements compared to baseline methods, further enhanced by the TRUE benchmark by applying an ensemble approach utilizing ANLI (standard NLI), SummaC (increased granularity NLI) as well as the Q² (QG-QA) method. The ensemble was shown to beat every standalone evaluation metric previously

discussed on all datasets benched in TRUE, emphasizing the need for combined, complementary evaluation metrics.

Recently, the evaluation paradigm has shifted to LLM-as-a-judge approaches. While early versions like GPTScore [34] tried evaluating based on model internals, using the conditional token probability, thus the confidence of a model in a given answer, they have quickly been shown to be much less effective than prompting the LLM directly to grade [35]. The performance of LLMs in evaluation tasks can additionally be enhanced by applying a Chain-of-thought approach, prompting the model to break down complex tasks into simpler ones and adding intermediate reasoning steps. Closely resembling human cognitive processes, this method yields significant improvements, especially with larger models and problems of increasing complexity [36], reaching agreement rates comparable to human judges for open-ended multi-turn dialogues [37], thus eliminating human annotation constraints.

When applied to long-form text containing multiple pieces of information, binary judgements become insufficient as they cannot reflect partial attribution. FactScore [38] addresses this by breaking down answers into atomic facts using an LLM, obtaining a percentage measure of supportedness. AttrScore [39] introduces additional evaluation criteria, such as measures for attribution, extrapolation and contradiction, in order to further characterize the type of attribution error. Similarly, CiteEval [40] introduces penalties for attributions containing unnecessary noise and redundant citations while advocating for evaluation against full source documents, in order to not lose information in the surrounding context. The RAGAS [41] framework operationalized these findings, providing a suite of automated metrics like faithfulness, answer relevance and context relevance, without requiring human annotation. However, the importance of penalizing undesired behaviors is further emphasized by the MT-Benchmark [37], showing that LLMs have a multitude of biases, such as favoring longer and LLM-generated text spans or switching preferences between answers when presented in a different order. Furthermore, LLMs perform remarkably poorly with fine-grained information, with a majority of mistakes being caused by overlooking numbers and entities or hallucinating logical connections [42]. Using increasingly large models additionally introduces high computational cost requirements. To address this, systems like ARES [43] rely on LLM-generated, domain-specific, synthetic training data to finetune smaller models. Using Prediction-Powered-Inference (PPI), ARES leverages a small human-annotated dataset to calculate the LLM’s annotation error on the synthetic data, the "rectifier", to provide statistical confidence bounds and enhance its evaluation, beating existing zero-shot solutions like RAGAS.

Automatic evaluation of cited text span attribution remains a challenging task. Despite the wide range of approaches, no single metric being capable of reliably and comprehensively assessing attribution performance as a whole has emerged. Increasingly complex tasks, such as multi-hop questions and challenging application domains like scientific QA, introduce further evaluation dimensions. Consequently, combining multiple evaluation metrics with complementary strengths becomes necessary to measure task-specific aspects within a unified framework.

3. Methodology and Approach

3.1. Dataset

To curate a dataset of high-quality, expert-level scientific question-answer pairs, a fully automated pipeline is applied. It consists of three consecutive steps: document selection, question-answer generation and quality assurance in order to ensure that generated questions and answers are scientifically rigorous and strictly dependent on multi-document information synthesis in order to be answered, while also reliably providing associated ground-truth documents alongside them.

3.1.1. Document selection

To match the underlying document corpus, selectable papers are restricted to those of the unarXive 2022 dataset, yet for other applications, this could be expanded at will. A randomly selected anchor paper as well as its cited references serve as the foundation for generating a single question-answer pair. For each paper, a document-level SPECTER [44] vector embedding is retrieved via the Semantic Scholar academic graph API [45]. Leveraging citation graphs, these vectors provide dense representations of a paper’s title and abstract, which are subsequently used to calculate the cosine similarity and thus the semantic similarity of each paper with the anchor document. To ensure topical alignment while excluding highly redundant work, a subset of candidate papers with a cosine similarity between 0.7 and 0.9 is selected. Four of these papers are then chosen at random to be used alongside the anchor paper for generation. Restricting the model input to five papers in total provides a multitude of topical overlaps for the model to choose from without resulting in exceedingly long prompts, mitigating the risk of "Lost in the Middle" effects. If a chosen anchor paper fails to provide at least four appropriate references, it is considered insufficient to be used in the dataset.

3.1.2. Question-Answer generation

While modern LLMs offer increasingly large context windows, usually even above 100,000 tokens, empirical studies have shown severe performance degradation in tasks like complex reasoning, data retrieval and logic problems, accompanied by drastically increased hallucination [46]. In the generation step, the pipeline applies a moonshotai/Kimi-K2.6 [47] model, supporting a context window of approximately 256,000 tokens. In practice, however, the "effective context window", in which models only suffer from minimal performance degradations with increasingly long inputs, has been shown by empirical studies [48] and practitioner reports to be drastically lower, with some studies estimating it as low as 40% of the maximum supported context window [49]. To ensure high reasoning capabilities, the pipeline applies an adaptive character limit for each paper text, enforcing a total upper bound of 200,000 characters for all papers combined, which conservatively translates to roughly 66,000 tokens for scientific paper texts [50], or about 26% of Kimi-K2-6’s supported

context size. Papers are assigned input space according to the following formula, allowing for full and efficient utilization of the context window, even if papers are highly variable in length:

$$C_{adaptive} = \left\lfloor C_{base} + \left(\frac{l_{paper}}{\sum_{p \in P_{above}} len(p)} \right) \times \sum_{s \in P_{below}} (C_{base} - len(s)) \right\rfloor \quad (3.1)$$

Paper texts longer than their adaptive character limit are then head- and tail-truncated, retaining the first and last $\lfloor \frac{C_{adaptive}}{2} \rfloor$ characters, which are the most information-dense parts of scientific papers.

Subsequently, the LLM is prompted using a two-step strategy. Firstly, the model is tasked with selecting the best pair of papers (A, B), where paper A provides an intermediate entity, method, dataset or a similar element, and paper B evaluates, extends, contradicts or relies on A’s intermediate element. Following that, the LLM is tasked with generating a single question-answer pair based on its selected pair, whose answer can only be derived by following a sequential reasoning process across both papers, applying paper A’s intermediate element to paper B. Additionally, the model is instructed to answer the question as precisely as possible, providing a precision-biased subset of the underlying papers used during generation alongside.

3.1.3. Quality assurance

In order to ensure the dataset’s integrity, two separate evaluation steps are applied, both using an openai/gpt-oss-120b [51] model. The first judge evaluates each Question-Answer pair along six distinct dimensions, providing a judgement on how well each criterion is satisfied on a three-point ordinal scale.

1. **Multi-Hop Validity** assesses whether answering the question requires integrating information from both papers. Thus, this criterion is only satisfied if neither paper provides sufficient information for deriving the complete answer. A question only satisfies this criterion if neither paper alone provides sufficient evidence to derive the complete answer.
2. **Dependency Strength** evaluates whether the steps in the reasoning chain are sequentially dependent in order for the conclusion to be drawn. Thus, the second reasoning step cannot be solved without first applying or interpreting the first one.
3. **Non-Decomposability** verifies that questions cannot be broken down into independent single-hop sub-questions, allowing for reasoning shortcuts or simply combining sub-question answers. Instead, reasoning steps must be connected in such a way that separately solving them would not be sufficient for deriving an answer.
4. **Evidence Distribution** examines whether both papers contribute distinct and necessary evidence, ensuring that required evidence is spread across both papers and that questions cannot be answered by using a single dominant source with redundant information.
5. **Answerability** assesses whether the answer is fully supported and obtainable by only relying on information present in the papers, without requiring external knowledge or speculative

inference not explicitly supported by either paper. Is the answer fully supported by the provided evidence?

6. **Decontextualization** validates that the question is self-contained and unambiguous in an open-domain setting. Thus, questions cannot contain context-specific references and must instead introduce relevant concepts, methods or entities by themselves in order to satisfy this criterion.

An analysis of disconnected reasoning by Trivedi et al. [52] demonstrated that many alleged multi-hop question-answering datasets do in fact not require true multi-hop reasoning, but can instead be solved through reasoning shortcuts, partial evidence or decomposition into independent sub-questions. Using a methodology inspired by the DIRE [52] framework, the second judge verifies that questions cannot be answered under incomplete evidence, namely a subset $P' \subset P$ of ground truth papers P , but instead require all ground truth documents for a complete reasoning chain and answer. Thus, the test outcome satisfies:

$$(\forall P' \subset P : \neg \text{Pass}(P')) \wedge \text{Pass}(P)$$

The final dataset is then obtained by filtering question-answer pairs to only those successfully passing both judges. In order to verify the pipeline and the judges, human verification will be applied on a subset.

3.2. Cited text span extraction

Constrained by the SQuAI system’s attribution granularity, which natively provides exactly one source document per answer sentence, the source span extraction task in this work is formulated as a post-hoc attribution step for providing one source span per answer sentence. The retrieval methods themselves, however, are applicable to other systems with alternative attribution granularities. This constitutes a limitation for this specific system, especially in cases where one attribution scope (e.g., a sentence) contains multiple claims requiring evidence, but also avoids additional ambiguity introduced by finer attribution granularity. For the purpose of comparing different extraction methods under identical conditions and constraints, however, this granularity is sufficient in order to measure whether extraction variants are capable of identifying relevant evidence, with results being directly comparable.

As previously discussed, cited text span attribution can be approached through a number of different paradigms, with methods being introduced and applied under vastly different circumstances. In the following, a hierarchy of increasingly complex post-hoc extraction methods that will be applied on a unified dataset is introduced. All methods except the native SQuAI span extraction and the generative LLM method receive tuples composed of an answer sentence and a set of eligible source spans $\mathcal{W}(P)$ from the paper P , annotated as the factual basis for that sentence during question answering. The set of source spans is composed of floating windows of contiguous sentences in the source document, up to a length of five, and can be denoted as:

$$\mathcal{W}(P) = \bigcup_{k=1}^5 \{s_i \oplus s_{i+1} \oplus \dots \oplus s_{i+k-1} \mid 1 \leq i \leq n - k + 1\},$$

where k specifies the window size and $\oplus$ denotes string concatenation with whitespace.

Extraction Method	Floating windows	Retrieval Category	Candidate Selection (Top 10)	Final Selection (Top 1)
1	×	Native	-	SQuAI
2	✓	Dense	-	Bi-encoder
3	✓	Sparse	-	BM25
4	✓	Hybrid	Bi-encoder	BM25
5	✓	Hybrid	BM25	Bi-encoder
6	✓	Hybrid	-	BM25 & Bi-encoder ²
7	✓	Cross-encoder	Bi-encoder	Cross-encoder
8	✓	Cross-encoder	BM25	Cross-encoder
9	✓	Cross-encoder	BM25 & Bi-encoder ²	Cross-encoder
10	×	Generative	-	LLM

Table 3.1.: Cited text span extraction methods overview

3.2.1. Native SQuAI extraction baseline

The first method relies on SQuAI’s native text span extraction mechanism, utilizing simple lexical matching. Since it is only capable of providing one span per paper, even if it is cited repeatedly in different contexts, it is limited by design. Nevertheless, representing the attribution behaviour of the SQuAI system, it serves as a native baseline for further extraction pipelines to compare against.

3.2.2. Dense retrieval

Applying a bi-encoder, this method selects the best-fitting source span out of all floating windows based on the similarity of the span’s dense vector representation with that of the answer sentence it is attributed to. It measures the suitability and sufficiency of relying exclusively on semantic similarity for source span attribution. As a dense retrieval baseline, it is capable of identifying appropriate source spans with different wording, but is prone to overlooking precise terminology.

3.2.3. Sparse retrieval

The third method utilizes BM25, reranking source spans based on their lexical similarity with the corresponding answer sentence. The highest-scoring span is then selected as the final attributed

² using Reciprocal Rank Fusion

span. Representing the sparse retrieval baseline, BM25 is especially effective with precise terminology, but less robust to paraphrasing. Its purpose is to test whether focusing on specific terms, entities or important keywords alone proves effective for cited text span extraction.

3.2.4. Hybrid retrieval

Hybrid methods aim to combine sparse and dense retrieval, leveraging their distinct characteristics. Thus, these variants evaluate whether the retrieval performance can be enhanced by combining both approaches into a single method. The fourth method applies a two-stage retrieval method, with the bi-encoder selecting the top ten candidate windows and BM25 selecting the final span, while method five reverses their order. While the former prioritizes semantic recall during candidate selection and lexical precision during final span selection, the latter reverses this order. Comparing both allows for analysis of whether either of those paradigms is better suited for candidate or final span selection. The sixth method combines both lexical and semantic matching into a single score using Reciprocal Rank Fusion (RRF). Instead of sequential application, this variant evaluates hybrid retrieval performance without an imposed order between the components themselves. Thus, this combination allows for assessment of whether the paradigms capture complementary aspects of the attribution task, resulting in performance enhancements when applied at the same time.

3.2.5. Cross-Encoder reranking

Overcoming the limitation of BM25 and bi-encoder retrieval processing the answer sentence and source span in isolation, the following methods introduce cross-encoders, allowing for joint processing of both. This promises improvements, especially for context windows where fine-grained semantic differences such as negations or conditions as well as important entities or numerical values determine whether the answer sentence is merely related, contradicted or truly supported. These variants therefore evaluate whether joint answer-span comparison translates into better evidence selection in the form of cited text spans. However, as cross-encoders are far more computationally expensive than previous methods and thus not feasible to run on the whole set of floating window candidate spans, these methods instead use different candidate selectors, with the cross-encoder selecting the best-fitting final span for the answer sentence. Methods seven and eight apply semantic and lexical matching, respectively, while method nine utilizes both together as a candidate selector using RRF. Comparing these variations allows assessing which candidate selection strategy provides the best basis for a cross-encoder and therefore improves overall attribution quality.

3.2.6. Generative LLM

The last method uses a generative LLM to extract the best supporting span from the source document. Unlike previous methods, it does not rerank candidate spans, but instead prompts the model to identify the best-suited text span directly. Representing a generative approach, it is used to test whether an LLM’s reasoning capabilities can improve evidence selection beyond retrieval and reranking, serving as a high-cost reference method.

3.3. System evaluation

The performance of the text span extraction methods can only be evaluated conditionally on the outputs of preceding pipeline steps, since post-hoc attribution is heavily dependent on their results. The evaluation is therefore split into assessing the suitability of the upstream inputs for text span extraction and evaluating the quality of the extracted spans themselves. This separation prevents text-span evaluation results from being confounded by errors introduced before span extraction.

3.3.1. Upstream system evaluation

Upstream inputs leading to bad span attributions can be flawed for a number of reasons, including:

1. Retrieval failure: The system retrieves or filters the wrong or insufficient subset of papers.
2. Generation failure: The system generates a factually wrong, overgeneralized, hallucinated or insufficiently supported answer.
3. Document attribution failure: The system fails to attribute or incorrectly attributes source documents to answer sentences.

These must be addressed separately when evaluating the attributed spans.

Retrieval failure can be analyzed by comparing the retrieved documents with those used during question generation, making it possible not only to evaluate the retrieval capabilities of the SQuAI system itself using precision, recall and F1-scores, but also the conditional performance of attribution methods given only unrelated, partial or complete evidence sources.

Generation failure is examined by comparing the system-generated answer with the ground-truth answer provided in the dataset. This comparison evaluates whether the system’s answer covers at least the factual content of the reference, including all relevant facts and without introducing contradictory information. Exact lexical overlap is not required, as the comparison is performed by an LLM-as-a-judge, whose reliability is validated through human annotation on a small subset.

Document attribution failure can be evaluated by running extraction methods on all retrieved source documents, regardless of the sentence-specific document annotation by the answering LLM. This allows for comparing the best spans from both the annotated document and the whole set of retrieved papers. If no sufficient span can be extracted from the former, while the latter provides one, the failure indicates incorrect document attribution rather than a failure of the span extraction method itself.

3.3.2. Span evaluation

The attributed spans are evaluated across multiple distinct dimensions in span-sentence pairs, with the goal of assessing whether they provide sufficient evidence for claims made in their respective answer sentence. This evaluation applies both an LLM-as-a-judge approach with RAGAS-inspired metrics, such as faithfulness, answer relevance and context relevance [41], as well as an NLI model to check for logical entailment between the span and the answer sentence. Additionally, the runtime required to extract spans will be compared across methods, allowing the evaluation to capture the trade-off between attribution quality and computational cost.

4. Timetable

The development of this thesis is based on extensive literature research conducted prior to this Exposé in order to properly outline the requirement and development scope of the dataset, the system architecture as well as the thesis itself. The initial setup on the High-Performance Computing Cluster has been completed. Currently, the literature search is nearing completion and the curation of the dataset is in progress, resulting in the following timeline for thesis completion:

Stage	Timeframe
Literature Search	25/05/2026 – 31/05/2026
System design	01/06/2026 – 14/06/2026
Implementation	15/06/2026 – 28/06/2026
Evaluation	29/06/2026 – 12/07/2026
Conclusion	13/07/2026 – 02/08/2026

Bibliography

- [1] Tobias Schreieder, Tim Schopf, and Michael Färber. *Attribution, Citation, and Quotation: A Survey of Evidence-based Text Generation with Large Language Models*. 2025. arXiv: 2508.15396 [cs.CL]. URL: <https://arxiv.org/abs/2508.15396>.
- [2] Jie Huang and Kevin Chang. “Citation: A Key to Building Responsible and Accountable Large Language Models”. In: *Findings of the Association for Computational Linguistics: NAACL 2024*. Ed. by Kevin Duh, Helena Gomez, and Steven Bethard. Mexico City, Mexico: Association for Computational Linguistics, June 2024, pp. 464–473. DOI: 10.18653/v1/2024.findings-naacl.31. URL: <https://aclanthology.org/2024.findings-naacl.31/>.
- [3] Ines Besrouer et al. “SQuAI: Scientific Question-Answering with Multi-Agent Retrieval-Augmented Generation”. In: *Proceedings of the 34th ACM International Conference on Information and Knowledge Management*. CIKM ’25. ACM, Nov. 2025, pp. 6603–6608. DOI: 10.1145/3746252.3761471. URL: <http://dx.doi.org/10.1145/3746252.3761471>.
- [4] Tarek Saier, Johan Krause, and Michael Färber. “unarXive 2022: All arXiv Publications Pre-Processed for NLP, Including Structured Full-Text and Citation Network”. In: *2023 ACM/IEEE Joint Conference on Digital Libraries (JCDL)*. IEEE, 2023, pp. 66–70. DOI: 10.1109/jcdl57899.2023.00020. URL: <http://dx.doi.org/10.1109/JCDL57899.2023.00020>.
- [5] Pradeep Dasigi et al. *A Dataset of Information-Seeking Questions and Answers Anchored in Research Papers*. 2021. arXiv: 2105.03011 [cs.CL]. URL: <https://arxiv.org/abs/2105.03011>.
- [6] Tianyu Gao et al. *Enabling Large Language Models to Generate Text with Citations*. 2023. arXiv: 2305.14627 [cs.CL]. URL: <https://arxiv.org/abs/2305.14627>.
- [7] Shraman Pramanick, Rama Chellappa, and Subhashini Venugopalan. *SPIQA: A Dataset for Multimodal Question Answering on Scientific Papers*. 2025. arXiv: 2407.09413 [cs.CL]. URL: <https://arxiv.org/abs/2407.09413>.
- [8] Yoonjoo Lee et al. “QASA: Advanced Question Answering on Scientific Articles”. In: *Proceedings of the 40th International Conference on Machine Learning*. Ed. by Andreas Krause et al. Vol. 202. Proceedings of Machine Learning Research. PMLR, July 2023, pp. 19036–19052. URL: <https://proceedings.mlr.press/v202/lee23n.html>.
- [9] Chaitanya Malaviya et al. *ExpertQA: Expert-Curated Questions and Attributed Answers*. 2024. arXiv: 2309.07852 [cs.CL]. URL: <https://arxiv.org/abs/2309.07852>.
- [10] Shruti Singh, Nandan Sarkar, and Arman Cohan. “SciDQA: A Deep Reading Comprehension Dataset over Scientific Papers”. In: *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*. Ed. by Yaser Al-Onaizan, Mohit Bansal, and Yun-Nung Chen. Miami, Florida, USA: Association for Computational Linguistics, Nov. 2024, pp. 20908–20923. DOI: 10.18653/v1/2024.emnlp-main.1163. URL: <https://aclanthology.org/2024.emnlp-main.1163/>.

- [11] Tim Baumgärtner, Ted Briscoe, and Iryna Gurevych. *PeerQA: A Scientific Question Answering Dataset from Peer Reviews*. 2025. arXiv: 2502.13668 [cs.CL]. URL: <https://arxiv.org/abs/2502.13668>.
- [12] Yash Saxena et al. *Generation-Time vs. Post-hoc Citation: A Holistic Evaluation of LLM Attribution*. 2025. arXiv: 2509.21557 [cs.CL]. URL: <https://arxiv.org/abs/2509.21557>.
- [13] Gorjan Radevski et al. *On Synthesizing Data for Context Attribution in Question Answering*. 2025. arXiv: 2504.05317 [cs.IR]. URL: <https://arxiv.org/abs/2504.05317>.
- [14] Ori Ernst et al. “Where Did That Come From? Sentence-Level Error-Tolerant Attribution”. In: *Findings of the Association for Computational Linguistics: EMNLP 2025*. Ed. by Christos Christodoulopoulos et al. Suzhou, China: Association for Computational Linguistics, Nov. 2025, pp. 6400–6417. ISBN: 979-8-89176-335-7. DOI: 10.18653/v1/2025.findings-emnlp.339. URL: <https://aclanthology.org/2025.findings-emnlp.339/>.
- [15] Ansel MacLaughlin et al. *Context-Based Quotation Recommendation*. 2020. arXiv: 2005.08319 [cs.CL]. URL: <https://arxiv.org/abs/2005.08319>.
- [16] Sriram Balasubramanian et al. *Decomposition-Enhanced Training for Post-Hoc Attributions In Language Models*. 2025. arXiv: 2510.25766 [cs.CL]. URL: <https://arxiv.org/abs/2510.25766>.
- [17] Masana Khanyi, Milandre Lill, and Marcel Dunaiski. “Citation Context Extraction Using a Transformer-Based Model with a Question-Answering Task”. In: July 2025, pp. 293–306. ISBN: 978-3-031-96261-5. DOI: 10.1007/978-3-031-96262-2_19.
- [18] Eran Hirsch et al. *LAQuer: Localized Attribution Queries in Content-grounded Generation*. 2025. arXiv: 2506.01187 [cs.CL]. URL: <https://arxiv.org/abs/2506.01187>.
- [19] Anirudh Phukan et al. *Peering into the Mind of Language Models: An Approach for Attribution in Contextual Question Answering*. 2024. arXiv: 2405.17980 [cs.CL]. URL: <https://arxiv.org/abs/2405.17980>.
- [20] Samyadeep Basu et al. *On Mechanistic Circuits for Extractive Question-Answering*. 2025. arXiv: 2502.08059 [cs.CL]. URL: <https://arxiv.org/abs/2502.08059>.
- [21] Michele Bevilacqua et al. *Autoregressive Search Engines: Generating Substrings as Document Identifiers*. 2022. arXiv: 2204.10628 [cs.CL]. URL: <https://arxiv.org/abs/2204.10628>.
- [22] Jiangui Chen et al. “GERE: Generative Evidence Retrieval for Fact Verification”. In: *Proceedings of the 45th International ACM SIGIR Conference on Research and Development in Information Retrieval*. SIGIR ’22. ACM, July 2022, pp. 2184–2189. DOI: 10.1145/3477495.3531827. URL: <http://dx.doi.org/10.1145/3477495.3531827>.
- [23] Liwen Zheng et al. “MRR-FV: Unlocking Complex Fact Verification with Multi-Hop Retrieval and Reasoning”. In: *Proceedings of the AAAI Conference on Artificial Intelligence* 39.24 (Apr. 2025), pp. 26066–26074. DOI: 10.1609/aaai.v39i24.34802. URL: <https://ojs.aaai.org/index.php/AAAI/article/view/34802>.
- [24] Kishore Papineni et al. “Bleu: a Method for Automatic Evaluation of Machine Translation”. In: *Annual Meeting of the Association for Computational Linguistics*. 2002. URL: <https://api.semanticscholar.org/CorpusID:11080756>.

- [25] Chin-Yew Lin. “ROUGE: A Package for Automatic Evaluation of Summaries”. In: *Text Summarization Branches Out*. Barcelona, Spain: Association for Computational Linguistics, July 2004, pp. 74–81. URL: <https://aclanthology.org/W04-1013/>.
- [26] Tianyi Zhang et al. *BERTScore: Evaluating Text Generation with BERT*. 2020. arXiv: 1904.09675 [cs.CL]. URL: <https://arxiv.org/abs/1904.09675>.
- [27] Wei Zhao et al. *MoverScore: Text Generation Evaluating with Contextualized Embeddings and Earth Mover Distance*. 2019. arXiv: 1909.02622 [cs.CL]. URL: <https://arxiv.org/abs/1909.02622>.
- [28] Weizhe Yuan, Graham Neubig, and Pengfei Liu. *BARTScore: Evaluating Generated Text as Text Generation*. 2021. arXiv: 2106.11520 [cs.CL]. URL: <https://arxiv.org/abs/2106.11520>.
- [29] Hannah Rashkin et al. *Measuring Attribution in Natural Language Generation Models*. 2022. arXiv: 2112.12870 [cs.CL]. URL: <https://arxiv.org/abs/2112.12870>.
- [30] Or Honovich et al. *TRUE: Re-evaluating Factual Consistency Evaluation*. 2022. arXiv: 2204.04991 [cs.CL]. URL: <https://arxiv.org/abs/2204.04991>.
- [31] Wojciech Kryściński et al. *Evaluating the Factual Consistency of Abstractive Text Summarization*. 2019. arXiv: 1910.12840 [cs.CL]. URL: <https://arxiv.org/abs/1910.12840>.
- [32] Philippe Laban et al. *SummaC: Re-Visiting NLI-based Models for Inconsistency Detection in Summarization*. 2021. arXiv: 2111.09525 [cs.CL]. URL: <https://arxiv.org/abs/2111.09525>.
- [33] Or Honovich et al. “Q²: Evaluating Factual Consistency in Knowledge-Grounded Dialogues via Question Generation and Question Answering”. In: *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*. Ed. by Marie-Francine Moens et al. Online and Punta Cana, Dominican Republic: Association for Computational Linguistics, Nov. 2021, pp. 7856–7870. DOI: 10.18653/v1/2021.emnlp-main.619. URL: <https://aclanthology.org/2021.emnlp-main.619/>.
- [34] Jinlan Fu et al. *GPTScore: Evaluate as You Desire*. 2023. arXiv: 2302.04166 [cs.CL]. URL: <https://arxiv.org/abs/2302.04166>.
- [35] Yang Liu et al. *G-Eval: NLG Evaluation using GPT-4 with Better Human Alignment*. 2023. arXiv: 2303.16634 [cs.CL]. URL: <https://arxiv.org/abs/2303.16634>.
- [36] Jason Wei et al. *Chain-of-Thought Prompting Elicits Reasoning in Large Language Models*. 2023. arXiv: 2201.11903 [cs.CL]. URL: <https://arxiv.org/abs/2201.11903>.
- [37] Lianmin Zheng et al. *Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena*. 2023. arXiv: 2306.05685 [cs.CL]. URL: <https://arxiv.org/abs/2306.05685>.
- [38] Sewon Min et al. *FActScore: Fine-grained Atomic Evaluation of Factual Precision in Long Form Text Generation*. 2023. arXiv: 2305.14251 [cs.CL]. URL: <https://arxiv.org/abs/2305.14251>.
- [39] Xiang Yue et al. *Automatic Evaluation of Attribution by Large Language Models*. 2023. arXiv: 2305.06311 [cs.CL]. URL: <https://arxiv.org/abs/2305.06311>.

- [40] Yumo Xu et al. “CiteEval: Principle-Driven Citation Evaluation for Source Attribution”. In: *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. Ed. by Wanxiang Che et al. Vienna, Austria: Association for Computational Linguistics, July 2025, pp. 32759–32778. ISBN: 979-8-89176-251-0. DOI: 10.18653/v1/2025.acl-long.1574. URL: <https://aclanthology.org/2025.acl-long.1574/>.
- [41] Shahul Es et al. *Ragas: Automated Evaluation of Retrieval Augmented Generation*. 2025. arXiv: 2309.15217 [cs.CL]. URL: <https://arxiv.org/abs/2309.15217>.
- [42] Yifei Li et al. *AttributionBench: How Hard is Automatic Attribution Evaluation?* 2024. arXiv: 2402.15089 [cs.CL]. URL: <https://arxiv.org/abs/2402.15089>.
- [43] Jon Saad-Falcon et al. *ARES: An Automated Evaluation Framework for Retrieval-Augmented Generation Systems*. 2024. arXiv: 2311.09476 [cs.CL]. URL: <https://arxiv.org/abs/2311.09476>.
- [44] Arman Cohan et al. “SPECTER: Document-level Representation Learning using Citation-informed Transformers”. In: *ArXiv abs/2004.07180* (2020). URL: <https://api.semanticscholar.org/CorpusID:215768677>.
- [45] URL: <https://api.semanticscholar.org/api-docs/>.
- [46] Norman Paulsen. *Context Is What You Need: The Maximum Effective Context Window for Real World Limits of LLMs*. Sept. 2025. DOI: 10.48550/arXiv.2509.21361.
- [47] URL: <https://huggingface.co/moonshotai/Kimi-K2.6>.
- [48] Cheng-Ping Hsieh et al. *RULER: What’s the Real Context Size of Your Long-Context Language Models?* 2024. arXiv: 2404.06654 [cs.CL]. URL: <https://arxiv.org/abs/2404.06654>.
- [49] Weiwei Wang, Jiyong Min, and Weijie Zou. *Intelligence Degradation in Long-Context LLMs: Critical Threshold Determination via Natural Length Distribution Analysis*. 2026. arXiv: 2601.15300 [cs.CL]. URL: <https://arxiv.org/abs/2601.15300>.
- [50] Jonathan Roberts, Kai Han, and Samuel Albanie. *How Long Is a Piece of String? A Brief Empirical Analysis of Tokenizers*. 2026. arXiv: 2601.11518 [cs.CL]. URL: <https://arxiv.org/abs/2601.11518>.
- [51] URL: [openai/gpt-oss-120b](https://openai.com/gpt-oss-120b).
- [52] Harsh Trivedi et al. “Is Multihop QA in DiRe Condition? Measuring and Reducing Disconnected Reasoning”. In: *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*. Ed. by Bonnie Webber et al. Online: Association for Computational Linguistics, Nov. 2020, pp. 8846–8863. DOI: 10.18653/v1/2020.emnlp-main.712. URL: <https://aclanthology.org/2020.emnlp-main.712/>.